




pooled over 12 encoders, 10 tasks, 29 folds



one line per task (10 tasks)



predictions spread more than their correlation warrants



Medians over 52,200 encoder-task-fold-gene cells, from the stored columns by the identity
 $R^2 = 2r\rho - \rho^2 - b^2/s_y^2$, nothing refitted.
 $R^2 = 0$ is predicting the test fold's own gene mean.

Rung 3 uses the test fold's mean, so it is diagnostic rather than achievable: it measures a slide-level shift no training-mean intercept can track.

Rung 4 is median(r^2) = 0.0997; per cell, r^2 is the ceiling for any level-and-scale recalibration.