




decodable from frozen features, within one PRAD patient



the same session axis, in the expression it is meant to predict



how much of the signature is tissue composition



IDC drops 3.0%, not the 17-27% the other three cross-patient tasks give -- it is also the only task here with a same-donor replicate pair.

The dotted 0.71 is the majority-class baseline for the resolution probe: 5 slides in one class, 2 in the other. Its withdrawn first version scored exactly that.

Left and middle: spatial-block cross-validation on PCA-256 features, so a held-out patch's neighbours are not in training. Middle regresses eight nuclear-morphology covariates out of each feature, fitted on training spots only; bars are per encoder, the annotated figures are means over the three. Right: PRAD patient 2 only, variance summed over the 50 target genes, method-of-moments components -- the raw between-session estimate is at or below zero for 28 of 50 genes.